

# 人工智慧概論 HW4

4112064214 侯竣奇

June 22, 2025

## 1 Summary

本次 HW4 作業的要求是在不使用 PyTorch 及 TensorFlow 等深度學習框架的前提下，使用 Numpy 來訓練一個卷積神經網路 (Convolutional Neural Network, CNN) 來分類 MNIST 資料集。可以另外訓練 CIFAR-10 或 ImageNet 資料集作為加分。

我在訓練過程中遇到過巢狀迴圈執行過於緩慢的問題，改用 `img2col` 的方式來加速，也使用 `Cupy` 來進行 GPU 加速。

訓練 MNIST 可以達到 98.60% 的準確率，同時我也完成了各項性能計算（精確率、召回率）、繪制學習曲線、卷積核視覺化（加分項）等任務。

作為加分挑戰，我也將模型應用於更複雜的 CIFAR-10 彩色圖片資料集，在這個資料集上取得了 59.01% 的準確率，並同樣完成了相關的性能評估與視覺化。

其餘詳細的性能分析與解說在檔案最後 Conclusion 的部分有提及。

## 2 Step 0: 下載資料集

```
[ ]: # 延用 HW3 中下載 mnist 的程式
def get_mnist():
    """
    The code to download the mnist data original came from
    https://cntk.ai/pythondocs/CNTK_103A_MNIST_DataLoader.html
    """

    import gzip
    import numpy as np
    import os
    import struct

    from urllib.request import urlretrieve

    def load_data(src, num_samples):
        print("Downloading " + src)
        gzfname, h = urlretrieve(src, "./delete.me")
        print("Done.")
        try:
            with gzip.open(gzfname) as gz:
```

```

        n = struct.unpack("I", gz.read(4))
        # Read magic number.
        if n[0] != 0x30800000:
            raise Exception("Invalid file: unexpected magic_
↪number.")

        # Read number of entries.
        n = struct.unpack(">I", gz.read(4))[0]
        if n != num_samples:
            raise Exception(f"Invalid file: expected_
↪{num_samples} entries.")
        crow = struct.unpack(">I", gz.read(4))[0]
        ccol = struct.unpack(">I", gz.read(4))[0]
        if crow != 28 or ccol != 28:
            raise Exception("Invalid file: expected 28 rows/
↪cols per image.")
        # Read data.
        res = np.frombuffer(gz.read(num_samples * crow *
↪ccol), dtype=np.uint8)
        finally:
            os.remove(gzfname)
        # 這次我們不在這裡做 normalize，留到 CuPy 陣列處理
        return res.reshape((num_samples, crow, ccol))

def load_labels(src, num_samples):
    print("Downloading " + src)
    gzfname, h = urlretrieve(src, "./delete.me")
    print("Done.")
    try:
        with gzip.open(gzfname) as gz:
            n = struct.unpack("I", gz.read(4))
            # Read magic number.
            if n[0] != 0x10800000:
                raise Exception("Invalid file: unexpected magic_
↪number.")

            # Read number of entries.
            n = struct.unpack(">I", gz.read(4))
            if n[0] != num_samples:
                raise Exception(f"Invalid file: expected_
↪{num_samples} rows.")
            # Read labels.
            res = np.frombuffer(gz.read(num_samples), dtype=np.
↪uint8)
            finally:
                os.remove(gzfname)
            return res.reshape(num_samples)

```

```

def try_download(data_source, label_source, num_samples):
    data = load_data(data_source, num_samples)
    labels = load_labels(label_source, num_samples)
    return data, labels

server = "https://raw.githubusercontent.com/fgnt/mnist/master"

# URLs for the train image and label data
url_train_image = f"{server}/train-images-idx3-ubyte.gz"
url_train_labels = f"{server}/train-labels-idx1-ubyte.gz"
num_train_samples = 60000

print("Downloading train data")
train_features, train_labels = try_download(url_train_image,
url_train_labels, num_train_samples)

# URLs for the test image and label data
url_test_image = f"{server}/t10k-images-idx3-ubyte.gz"
url_test_labels = f"{server}/t10k-labels-idx1-ubyte.gz"
num_test_samples = 10000

print("Downloading test data")
test_features, test_labels = try_download(url_test_image,
url_test_labels, num_test_samples)

return train_features, train_labels, test_features, test_labels

```

```

[ ]: train_features_np, train_labels_np, test_features_np, test_labels_np =
get_mnist()

```

Downloading train data

Downloading https://raw.githubusercontent.com/fgnt/mnist/master/  
↳ train-images-

idx3-ubyte.gz

Done.

Downloading https://raw.githubusercontent.com/fgnt/mnist/master/  
↳ train-labels-

idx1-ubyte.gz

Done.

Downloading test data

Downloading https://raw.githubusercontent.com/fgnt/mnist/master/  
↳ t10k-images-

idx3-ubyte.gz

Done.

Downloading https://raw.githubusercontent.com/fgnt/mnist/master/  
↳ t10k-labels-

idx1-ubyte.gz

Done.

### 3 Step 1: 資料前處理 (Data Preprocessing)

資料下載完成後，接下來我們要進行一系列的處理，才能將其輸入到 CNN 模型中。

#### 3.1 Step 1-1: 讀取資料並轉換為 CuPy 陣列

在 HW4 中我們改用 CuPy 來加速原先 Numpy 的運算速度

```
[ ]: import cupy as cp
import numpy as np

# 使用 cp.asarray() 將資料從 CPU memory 移至 GPU memory
train_x_orig = cp.asarray(train_features_np)
train_y_orig = cp.asarray(train_labels_np)
test_x_orig = cp.asarray(test_features_np)
test_y_orig = cp.asarray(test_labels_np)

# --- 資料塑形與正規化 (Reshape & Normalize) ---
# CNN 需要的輸入格式為 (樣本數, 色彩頻道, 高度, 寬度)
# MNIST 是灰階圖片, 所以 Channel 為 1
# 我們也在此將像素值從 0-255 正規化到 0-1 之間
train_x = train_x_orig.reshape(train_x_orig.shape[0], 1, 28, 28) / 255.
test_x = test_x_orig.reshape(test_x_orig.shape[0], 1, 28, 28) / 255.0
```

```
[ ]: # 驗證一下 shape 和 type
print("--- Shapes of CuPy arrays ---")
print("train_x.shape:", train_x.shape)
print("train_y_orig.shape:", train_y_orig.shape)
print("test_x.shape:", test_x.shape)
print("test_y_orig.shape:", test_y_orig.shape)
print("\n--- Data types ---")
print("type of train_x:", type(train_x))
print("type of train_y_orig:", type(train_y_orig))
```

```
--- Shapes of CuPy arrays ---
train_x.shape: (60000, 1, 28, 28)
train_y_orig.shape: (60000,)
test_x.shape: (10000, 1, 28, 28)
test_y_orig.shape: (10000,)
```

```
--- Data types ---
type of train_x: <class 'cupy.ndarray'>
type of train_y_orig: <class 'cupy.ndarray'>
```

## 3.2 Step 1-2: 對 Label 進行 One-Hot Encode

```
[ ]: def one_hot_encode(labels, num_classes):
    # 創建一個全為 0 的矩陣，shape 為 (樣本數, 類別數)
    one_hot = np.zeros((labels.size, num_classes))
    # 在對應的類別位置上填上 1
    one_hot[np.arange(labels.size), labels] = 1
    return one_hot

[ ]: # 我們的類別是數字 0-9，所以有 10 個類別
num_classes = 10
train_y = one_hot_encode(train_y_orig, num_classes)
test_y = one_hot_encode(test_y_orig, num_classes)

print("One-Hot encode 前的 train_y.shape:", train_y_orig.shape)
print("One-Hot 後的 train_y.shape:", train_y.shape)

print(f"一個 One-Hot 編碼範例 (原始 label 是 {train_y_orig[0]}):_
    ↳{train_y[0]}")
```

One-Hot encode 前的 train\_y.shape: (60000,)

One-Hot 後的 train\_y.shape: (60000, 10)

一個 One-Hot 編碼範例 (原始 label 是 5): [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

## 4 Step 2: 實作 CNN

### 4.1 Step 2-2: 初始化權重

初始化 CNN 的權重與偏置

參數:

- input\_dim – 輸入圖片的維度 (C, H, W)
- n\_filters, filter\_size, conv\_stride, conv\_padding – 卷積層的 hyperparameter
- pool\_size, pool\_stride – 池化層的超參數
- n\_h, n\_y – 全連接層與輸出層的神經元數量

Return:

- parameters – 一個包含所有權重與 bias 的 Python dict

```
[ ]: def initialize_parameters(input_dim=(1, 28, 28),
                              n_filters=8, filter_size=3, conv_stride=1,
                              ↳conv_padding=1,
                              pool_size=2, pool_stride=2,
                              n_h=128, n_y=10):

    np.random.seed(42)

    C_in, H_in, W_in = input_dim
```

```

# 1. 初始化卷積層權重
W_conv = cp.random.randn(n_filters, C_in, filter_size,
↪filter_size) * 0.01
b_conv = cp.zeros((n_filters, 1))

# 2. 動態計算全連接層的輸入維度
# 計算卷積層的輸出維度
H_conv = int((H_in - filter_size + 2 * conv_padding) /
↪conv_stride) + 1
W_conv_out = int((W_in - filter_size + 2 * conv_padding) /
↪conv_stride) + 1

# 計算池化層的輸出維度 (池化層不補零)
H_pool = int((H_conv - pool_size) / pool_stride) + 1
W_pool_out = int((W_conv_out - pool_size) / pool_stride) + 1

# 扁平化後的維度
fc_input_dim = n_filters * H_pool * W_pool_out

# print(f" 根據設定計算，全連接層的輸入維度為: {fc_input_dim}")

# 3. 初始化全連接層與輸出層權重
W_fc = cp.random.randn(fc_input_dim, n_h) * 0.01
b_fc = cp.ones((1, n_h)) * 0.01
W_out = cp.random.randn(n_h, n_y) * 0.01
b_out = cp.zeros((1, n_y))

parameters = {
    "W_conv": W_conv, "b_conv": b_conv,
    "W_fc": W_fc, "b_fc": b_fc,
    "W_out": W_out, "b_out": b_out
}

return parameters

```

```

[ ]: # MNIST
# 輸入維度: (1, 28, 28), 卷積設定: 3x3 filter, stride 1, padding 1 -> 輸出 (28, 28)
# 池化設定: 2x2, stride 2 -> 輸出 (14, 14)

params = initialize_parameters(input_dim=(1, 28, 28), n_filters=8,
↪filter_size=3)

print("\n--- Parameter Shapes ---")
print("W_conv Shape: " + str(params["W_conv"].shape))
print("W_fc Shape: " + str(params["W_fc"].shape))

```

```
--- Parameter Shapes ---  
W_conv Shape: (8, 1, 3, 3)  
W_fc Shape: (1568, 128)
```

## 4.2 Step 2-2: 實作卷積層 Forward Propagation

**zero\_pad**: 對 CuPy 陣列進行補零

參數:

- **X** – CuPy 陣列, shape (N, C, H, W)
- **pad** – 整數, 代表要補的圈數

Return:

- **X\_pad** – 補零後的陣列, shape (N, C, H + 2 \* pad, W + 2 \* pad)

---

**conv\_forward**: CNN 的前向傳播

參數:

- **A\_prev** – 上一層的輸出 (輸入圖片), shape (N, C\_prev, H\_prev, W\_prev)
- **W** – Filter 權重, shape (n\_filters, C\_prev, f, f)
- **b** – bias, shape (n\_filters, 1)
- **stride** – 滑動步長
- **padding** – 補零圈數

Return:

- **Z** – 卷積層的輸出, shape (N, n\_filters, H\_out, W\_out)
- **cache** – 用於反向傳播的快取

```
[ ]: def zero_pad(X, pad):  
    # cp.pad 的用法和 np.pad 相同  
    # ((0, 0), (0, 0), (pad, pad), (pad, pad)) 分別對應 4 個維度的補零設定  
    X_pad = cp.pad(X, ((0, 0), (0, 0), (pad, pad), (pad, pad)),  
mode='constant', constant_values=0)  
    return X_pad
```

```
[ ]: def conv_forward(A_prev, W, b, stride=1, padding=1):  
    # 1. 從 W 的維度取得濾波器資訊  
    (n_filters, C_prev, f, f) = W.shape  
    # 從 A_prev 的維度取得輸入資訊  
    (N, C_prev, H_prev, W_prev) = A_prev.shape  
  
    # 2. 計算輸出特徵圖的維度  
    H_out = int((H_prev - f + 2 * padding) / stride) + 1  
    W_out = int((W_prev - f + 2 * padding) / stride) + 1
```

```

# 3. 初始化輸出 Z
Z = cp.zeros((N, n_filters, H_out, W_out))

# 4. 對輸入 A_prev 進行補零
A_prev_pad = zero_pad(A_prev, padding)

# 5. 進行卷積運算 (使用巢狀迴圈)
for n in range(N): # 遍歷每一張圖片樣本
    a_prev_pad = A_prev_pad[n, :, :, :]
    for c in range(n_filters): # 遍歷每一個 Filter
        for h in range(H_out): # 遍歷輸出的垂直方向
            for w in range(W_out): # 遍歷輸出的水平方向
                # 定位當前滑動窗口的垂直和水平起始位置
                vert_start = h * stride
                vert_end = vert_start + f
                horiz_start = w * stride
                horiz_end = horiz_start + f

                # 選取當前的滑動窗口
                a_slice_prev = a_prev_pad[:, vert_start:vert_end,
↪horiz_start:horiz_end]

                # 進行對應元素相乘後相加，再加上 bias
                Z[n, c, h, w] = cp.sum(a_slice_prev * W[c, :, :, :
↪]) + b[c]

# 儲存快取以供反向傳播使用
cache = (A_prev, W, b, stride, padding)

return Z, cache

```

#### 4.2.1 Step 2-2-2: 優化前向傳播

在執行過程我發現這種四層 for 迴圈的寫法效能實在太低，所以我們改用 im2col 的寫法來改進

```

[ ]: def get_im2col_indices(x_shape, field_height, field_width, padding=1,
↪stride=1):
    # First figure out what the output shape is
    N, C, H, W = x_shape
    assert (H + 2 * padding - field_height) % stride == 0
    assert (W + 2 * padding - field_width) % stride == 0
    out_height = (H + 2 * padding - field_height) // stride + 1
    out_width = (W + 2 * padding - field_width) // stride + 1

    i0 = cp.repeat(cp.arange(field_height), field_width)
    i0 = cp.tile(i0, C)

```



```

    i1 = stride * cp.repeat(cp.arange(out_height), out_width)
    j0 = cp.tile(cp.arange(field_width), field_height * C)
    j1 = stride * cp.tile(cp.arange(out_width), out_height)
    i = i0.reshape(-1, 1) + i1.reshape(1, -1)
    j = j0.reshape(-1, 1) + j1.reshape(1, -1)

    k = cp.repeat(cp.arange(C), field_height * field_width).
↳ reshape(-1, 1)

    return (k, i, j)

def im2col_gpu(x, field_height, field_width, padding=1, stride=1):
    """ An implementation of im2col on GPU """
    # Zero-pad the input
    p = padding
    x_padded = cp.pad(x, ((0, 0), (0, 0), (p, p), (p, p)),
↳ mode='constant')

    k, i, j = get_im2col_indices(x.shape, field_height, field_width,
↳ padding, stride)

    cols = x_padded[:, k, i, j]
    C = x.shape[1]
    cols = cols.transpose(1, 2, 0).reshape(field_height * field_width,
↳ * C, -1)
    return cols

```

```

[ ]: def conv_forward_fast(A_prev, W, b, stride=1, padding=1):

    # 從 A_prev 的維度取得輸入資訊
    N, C_prev, H_prev, W_prev = A_prev.shape

    # 從 W (權重) 的維度取得 Filter 資訊
    n_filters, _, f, _ = W.shape

    # 計算輸出維度
    H_out = (H_prev - f + 2 * padding) // stride + 1
    W_out = (W_prev - f + 2 * padding) // stride + 1

    # 使用 im2col 將輸入圖片轉換為矩陣
    A_col = im2col_gpu(A_prev, f, f, padding=padding, stride=stride)

    # 將 Filter 權重攤平成矩陣
    W_row = W.reshape(n_filters, -1)

    # 核心運算：一次矩陣乘法 + bias
    Z = W_row @ A_col + b

```

```

# 將結果塑形成最終的輸出維度
Z = Z.reshape(n_filters, H_out, W_out, N)
Z = Z.transpose(3, 0, 1, 2) # 調整維度順序為 (N, n_filters, H_out, W_out)

# 快取 A_col 供反向傳播使用
cache = (A_prev, W, b, stride, padding, A_col)

return Z, cache

```

### 4.3 Step 2-3: 實作 Activation Function (ReLU)

ReLU 的反向傳播

參數:

- dA – 流回來的梯度
- cache\_Z – 前向傳播時的 Z 值

Return:

- dZ – 對 Z 的梯度

```

[ ]: def relu(Z):
    return cp.maximum(0, Z)

def relu_backward(dA, cache_Z):
    dZ = cp.array(dA, copy=True)
    # 當 Z <= 0 時，梯度也為 0
    dZ[cache_Z <= 0] = 0
    return dZ

```

### 4.4 Step 2-4: 實作池化層 Forward Propagation

實現 Max Pooling 的前向傳播

參數:

- A\_prev – 上一層的輸出 (ReLU), shape (N, C, H\_prev, W\_prev)
- pool\_size – 池化窗口的大小
- stride – 滑動步長

Return:

- A\_out – 池化層的輸出, shape (N, C, H\_out, W\_out)
- cache – 用於反向傳播的快取，包含 A\_prev 和超參數

```

[ ]: def max_pool_forward(A_prev, pool_size=2, stride=2):
    """
    """

```

```

# 取得輸入維度
(N, C, H_prev, W_prev) = A_prev.shape

# 計算輸出維度
H_out = int((H_prev - pool_size) / stride) + 1
W_out = int((W_prev - pool_size) / stride) + 1

# 初始化輸出
A_out = np.zeros((N, C, H_out, W_out))

# 進行最大池化運算
for n in range(N):
    for c in range(C):
        for h in range(H_out):
            for w in range(W_out):
                # 定位當前滑動窗口
                vert_start = h * stride
                vert_end = vert_start + pool_size
                horiz_start = w * stride
                horiz_end = horiz_start + pool_size

                # 選取滑動窗口
                a_slice = A_prev[n, c, vert_start:vert_end,
                                horiz_start:horiz_end]

                # 在窗口中找到最大值並賦給輸出
                A_out[n, c, h, w] = np.max(a_slice)

cache = (A_prev, (pool_size, stride))
return A_out, cache

```

#### 4.4.1 Step 2-4-2: 優化池化層前向傳播

```

[ ]: def max_pool_forward_fast(A_prev, pool_size=2, stride=2):

    N, C, H_prev, W_prev = A_prev.shape

    # 計算輸出維度
    H_out = 1 + (H_prev - pool_size) // stride
    W_out = 1 + (W_prev - pool_size) // stride

    # 透過 reshape 和 transpose 將每個池化窗口的元素分組
    A_reshaped = A_prev.reshape(N, C, H_out, pool_size, W_out,
                                pool_size)
    A_transposed = A_reshaped.transpose(0, 1, 2, 4, 3, 5)

```

```
# 在每個窗口內取最大值
A_out = cp.max(A_transposed, axis=(4, 5))

cache = (A_prev, A_out, (pool_size, stride))
return A_out, cache
```

## 4.5 Step 2-5: 完整的 Forward Propagation

流程如下：

1. CONV → RELU：對圖片進行卷積運算，然後傳給 ReLU
2. POOL: 對 ReLU 的輸出進行 Max Pooling
3. FLATTEN: 將池化後的特徵圖攤平成一維向量
4. FC→RELU: 將攤平的向量輸入全連接層，再通過一次 ReLU
5. FC→SOFTMAX: 最後通過輸出層，並使用 Softmax 函數得到最終的分類機率

Softmax 激活函數 (CuPy 版本)

為了數值穩定性，先減去 Z 中的最大值

```
[ ]: def softmax(Z):
    # Z.shape: (N, 10), N 是樣本數
    exp_scores = cp.exp(Z - cp.max(Z, axis=1, keepdims=True))
    return exp_scores / cp.sum(exp_scores, axis=1, keepdims=True)
```

實現 CNN 的完整前向傳播

參數:

- X – 輸入資料, shape (N, C, H, W)
- parameters – 包含所有權重的字典

Return:

- A\_out – 模型的最終輸出 (Softmax 的結果)
- cache – 包含所有中間層計算結果的字典，供反向傳播使用

```
[ ]: def forward_propagation(X, parameters):
    # 1. 取出權重
    W_conv = parameters['W_conv']
    b_conv = parameters['b_conv']
    W_fc = parameters['W_fc']
    b_fc = parameters['b_fc']
    W_out = parameters['W_out']
    b_out = parameters['b_out']

    # 2. CONV -> RELU -> POOL
    # 卷積層
    # 這裡我們使用固定的 stride=1, padding=1, pool_size=2, pool_stride=2
    Z_conv, cache_conv = conv_forward_fast(X, W_conv, b_conv,
    ↪ stride=1, padding=1)
```

```

# ReLU 激活
A_conv = relu(Z_conv)

# 池化層
A_pool, cache_pool = max_pool_forward_fast(A_conv, pool_size=2,
stride=2)

# 3. FLATTEN -> FC -> RELU
# 扁平化
N = A_pool.shape[0]
A_flat = A_pool.reshape(N, -1) # -1 會讓 cupy 自動計算維度

# 全連接層
Z_fc = cp.dot(A_flat, W_fc) + b_fc

# ReLU 激活
A_fc = relu(Z_fc)

# 4. OUTPUT LAYER -> SOFTMAX
Z_out = cp.dot(A_fc, W_out) + b_out
A_out = softmax(Z_out)

# 5. 儲存所有快取
# 這是反向傳播時的關鍵
cache = {
    "cache_conv": cache_conv, # (A_prev, W, b, stride, padding)
    "Z_conv": Z_conv, # Conv 輸出
    "A_conv": A_conv, # Conv 的 ReLU 輸出
    "cache_pool": cache_pool, # (A_conv, (pool_size, stride))
    "A_pool": A_pool, # Pool 輸出
    "A_flat": A_flat, # Flatten 輸出
    "Z_fc": Z_fc, # FC 輸出
    "A_fc": A_fc, # FC 的 ReLU 輸出
    "W_fc": W_fc, # FC 權重
    "W_out": W_out # Out 權重
}

return A_out, cache

```

## 4.6 Step 2-6: 損失函數 Loss Function

跟 HW3 一樣，使用 Cross Entropy

參數：

- **A\_out** – 模型的預測機率輸出, shape (N, num\_classes)
- **Y** – 真實的 One-Hot 標籤, shape (N, num\_classes)

Return:

- **loss** – 交叉熵損失值 (一個純量)

```
[ ]: def compute_loss(A_out, Y):  
  
    N = Y.shape[0]  
  
    # 加上一個極小值 1e-9 是為了避免 log(0) 的情況  
    loss = -cp.sum(Y * cp.log(A_out + 1e-9)) / N  
  
    return loss
```

## 4.7 Step 2-7: Backward Propagation

路徑如下：

$dLoss/dZ_{out} \rightarrow dLoss/dA_{fc} \rightarrow dLoss/dZ_{fc} \rightarrow dLoss/dA_{pool} \rightarrow dLoss/dA_{conv} \rightarrow dLoss/dZ_{conv}$

每一步，我們都會計算出對應層的權重梯度 ( $dW, db$ ) 以及傳遞給前一層的梯度 ( $dA$ )。

對於「Softmax + Cross Entropy」這個經典組合，起始梯度有個很簡潔的結果：

$$\frac{\partial Loss}{\partial_{out}} = A_{out} - Y$$

其中  $A_{out}$  是 Softmax 的預測機率， $Y$  是正解的 One-Hot 標籤。這大大簡化了我們的計算。我們就從這個梯度  $dZ_{out}$  開始。

參數:

- **A\_out** – 模型的預測輸出
- **Y** – 真實的 One-Hot 標籤
- **cache** – 前向傳播儲存的所有中間值
- **parameters** – 包含所有權重的字典

Return:

- **grads** – 包含所有權重與 bias 梯度的字典

```
[ ]: def backward_propagation(A_out, Y, cache, parameters):  
    grads = {}  
    N = Y.shape[0] # 批次大小  
  
    # 1. 輸出層的反向傳播  
    # 起始梯度 dZ_out  
    dZ_out = A_out - Y  
  
    # 對 W_out, b_out 的梯度  
    A_fc = cache['A_fc']  
    grads['dW_out'] = (1/N) * cp.dot(A_fc.T, dZ_out)
```

```

grads['db_out'] = (1/N) * cp.sum(dZ_out, axis=0, keepdims=True)

# 傳遞給前一層 (FC 層) 的梯度 dA_fc
W_out = parameters['W_out']
dA_fc = cp.dot(dZ_out, W_out.T)

# 2. 全連接層 (FC) 的反向傳播 ---
# 首先通過 ReLU 的反向傳播
Z_fc = cache['Z_fc']
dZ_fc = relu_backward(dA_fc, Z_fc) # 我們之前寫好的 relu_backward

# 對 W_fc, b_fc 的梯度
A_flat = cache['A_flat']
grads['dW_fc'] = (1/N) * cp.dot(A_flat.T, dZ_fc)
grads['db_fc'] = (1/N) * cp.sum(dZ_fc, axis=0, keepdims=True)

# 傳遞給前一層 (扁平層) 的梯度 dA_flat
W_fc = parameters['W_fc']
dA_flat = cp.dot(dZ_fc, W_fc.T)

# 3. 扁平層 (Flatten) 的反向傳播
# 這一步只是將梯度從 1D 向量恢復成 4D 陣列
A_pool = cache['A_pool']
dA_pool = dA_flat.reshape(A_pool.shape)

# 接下來的兩步 (池化層和卷積層) 比較複雜，我們先把函式框架寫好

# 4. 池化層 (Pooling) 的反向傳播
cache_pool = cache['cache_pool']
dA_conv = max_pool_backward_fast(dA_pool, cache_pool)

# 5. 卷積層 (Conv) 的反向傳播
# 首先通過 ReLU 的反向傳播
Z_conv = cache['Z_conv']
dZ_conv = relu_backward(dA_conv, Z_conv)

# 傳入卷積層反向傳播函式
cache_conv = cache['cache_conv']
dA_prev, dW_conv, db_conv = conv_backward_fast(dZ_conv, cache_conv)

grads['dW_conv'] = dW_conv
grads['db_conv'] = db_conv

return grads

```

上面的程式碼還無法執行，因為我們還缺少 `max_pool_backward` 和 `conv_backward` 這兩個最關鍵的函式。

#### 4.7.1 Step 2-7-2: 池化層的反向傳播

實現最大池化層的反向傳播

參數:

- **dA** – 池化層輸出的梯度, shape (N, C, H\_out, W\_out)
- **cache** – 前向傳播的快取 (A\_prev, (pool\_size, stride))

Return:

- **dA\_prev** – 傳遞給前一層的梯度

```
[ ]: def max_pool_backward(dA, cache):

    A_prev, (pool_size, stride) = cache
    (N, C, H_prev, W_prev) = A_prev.shape
    (N, C, H_out, W_out) = dA.shape

    dA_prev = cp.zeros_like(A_prev)

    for n in range(N):
        for c in range(C):
            for h in range(H_out):
                for w in range(W_out):
                    vert_start = h * stride
                    vert_end = vert_start + pool_size
                    horiz_start = w * stride
                    horiz_end = horiz_start + pool_size

                    # 選取前向傳播時的同一個窗口
                    a_slice = A_prev[n, c, vert_start:vert_end,
horiz_start:horiz_end]

                    # 找到最大值的位置 (mask)
                    mask = (a_slice == cp.max(a_slice))

                    # 將 dA 的梯度加到這個最大值的位置上
                    dA_prev[n, c, vert_start:vert_end, horiz_start:
horiz_end] += mask * dA[n, c, h, w]

    return dA_prev
```

#### 4.7.2 Step 2-7-3: 卷積層的反向傳播

實現卷積層的反向傳播

參數:

- **dZ** – 卷積層輸出的梯度, shape (N, n\_filters, H\_out, W\_out)
- **cache** – 前向傳播的快取 (A\_prev, W, b, stride, padding)



Return:

- dA\_prev, dW, db

```
[ ]: def conv_backward(dZ, cache):
    A_prev, W, b, stride, padding = cache
    (N, C_prev, H_prev, W_prev) = A_prev.shape
    (n_filters, C_prev, f, f) = W.shape
    (N, n_filters, H_out, W_out) = dZ.shape

    # 初始化梯度
    dA_prev = cp.zeros_like(A_prev)
    dW = cp.zeros_like(W)
    db = cp.zeros((n_filters, 1))

    # 對輸入進行補零 (為了計算 dA_prev)
    A_prev_pad = zero_pad(A_prev, padding)
    dA_prev_pad = zero_pad(dA_prev, padding)

    for n in range(N):
        # 選取單一樣本
        a_prev_pad = A_prev_pad[n]
        da_prev_pad = dA_prev_pad[n]

        for c in range(n_filters):
            for h in range(H_out):
                for w in range(W_out):
                    vert_start = h * stride
                    vert_end = vert_start + f
                    horiz_start = w * stride
                    horiz_end = horiz_start + f

                    # 定位滑動窗口
                    a_slice = a_prev_pad[:, vert_start:vert_end,
horiz_start:horiz_end]

                    # 計算梯度
                    # 1. 更新傳遞給前一層的梯度 dA_prev
                    da_prev_pad[:, vert_start:vert_end, horiz_start:
horiz_end] += W[c, :, :, :] * dZ[n, c, h, w]
                    # 2. 更新對濾波器的梯度 dW
                    dW[c, :, :, :] += a_slice * dZ[n, c, h, w]
                    # 3. 更新對偏置的梯度 db (此處先累加)
                    db[c] += dZ[n, c, h, w]

    # 去掉 dA_prev_pad 的補零部分
    if padding > 0:
```

```

        dA_prev[n, :, :, :] = da_prev_pad[:, padding:-padding,
padding:-padding]
    else:
        dA_prev[n, :, :, :] = da_prev_pad

    # 對 db 取平均
    db = db / N

    return dA_prev, dW, db

```

#### 4.7.3 Step 2-7-4: 優化卷積層反向傳播

跟前向傳播時一樣，執行時發現巢狀迴圈效率過差

```

[ ]: from cupy.lib.stride_tricks import as_strided

def col2im_gpu(cols, x_shape, field_height=3, field_width=3,
padding=1, stride=1):
    """ An implementation of col2im on GPU """
    N, C, H, W = x_shape
    H_padded, W_padded = H + 2 * padding, W + 2 * padding
    x_padded = cp.zeros((N, C, H_padded, W_padded), dtype=cols.dtype)
    k, i, j = get_im2col_indices(x_shape, field_height, field_width,
padding, stride)
    cols_resaped = cols.reshape(C * field_height * field_width, -1, N)
    cols_resaped = cols_resaped.transpose(2, 0, 1)

    # 使用 cp.add.at 這個原子操作來高效地將梯度加總回去
    cp.add.at(x_padded, (slice(None), k, i, j), cols_resaped)

    if padding == 0:
        return x_padded
    return x_padded[:, :, padding:-padding, padding:-padding]

```

```

[ ]: def conv_backward_fast(dZ, cache):
    """
    使用 col2im 實現的高效版卷積反向傳播
    """
    # 解開 cache
    A_prev, W, b, stride, padding, A_col = cache
    N, C, H, W_shape = A_prev.shape
    n_filters, _, _, _ = W_shape

    # 計算 db 的梯度
    db = cp.sum(dZ, axis=(0, 2, 3))
    db = db.reshape(n_filters, -1)

```

```

# 重新塑形 dZ 以便進行矩陣乘法
dZ_reshaped = dZ.transpose(1, 2, 3, 0).reshape(n_filters, -1)

# 計算 dW 的梯度
dW = dZ_reshaped @ A_col.T
dW = dW.reshape(W.shape)

# 計算傳遞給前一層的梯度 dA_prev
W_row = W.reshape(n_filters, -1)
dA_col = W_row.T @ dZ_reshaped
dA_prev = col2im_gpu(dA_col, A_prev.shape, W.shape[2], W.shape[3],
padding, stride)

return dA_prev, dW, db

```

#### 4.7.4 Step 2-7-5: 優化池化層的反向傳播

```

[ ]: def max_pool_backward_fast(dA, cache):
    # 從快取中取出前向傳播的輸入與輸出
    A_prev, A_out, (pool_size, stride) = cache

    # 1. 將池化層輸出的梯度 dA 上採樣
    dA_upsampled = dA.repeat(pool_size, axis=2).repeat(pool_size,
axis=3)

    # 2. 同樣地，將前向傳播的輸出 A_out 上採樣
    A_out_upsampled = A_out.repeat(pool_size, axis=2).
repeat(pool_size, axis=3)

    # 3. 建立遮罩：A_prev 中等於上採樣後 A_out 的位置，就是當時的最大值
    # 這些位置的值為 True (1)，其餘為 False (0)
    mask = (A_prev == A_out_upsampled)

    # 4. 將上採樣的梯度應用於遮罩，梯度只會流向最大值所在的位置
    dA_prev = dA_upsampled * mask

    return dA_prev

```

#### 4.8 Step 2-8: 更新權重

使用梯度下降法更新模型的參數

參數：

- **parameters** – 包含當前權重的字典
- **grads** – 包含梯度的字典
- **learning\_rate** – 學習率 (alpha)

Return:

- parameters – 更新後的參數字典

```
[ ]: def update_parameters(parameters, grads, learning_rate):

    # 根據梯度下降規則更新每一個參數
    parameters['W_conv'] -= learning_rate * grads['dW_conv']
    parameters['b_conv'] -= learning_rate * grads['db_conv']
    parameters['W_fc'] -= learning_rate * grads['dW_fc']
    parameters['b_fc'] -= learning_rate * grads['db_fc']
    parameters['W_out'] -= learning_rate * grads['dW_out']
    parameters['b_out'] -= learning_rate * grads['db_out']

    return parameters
```

## 5 Step 3: 整合出 CNN Model

### 5.1 Step 3-1: Create the CNN Model

完整的 CNN 訓練模型

Return:

- trained\_parameters – 訓練完成後的模型參數
- history – 包含每個 epoch 的 cost 和 accuracy 的字典

```
[ ]: import time # 引入 time 模組來計算訓練時間
import copy # 我們需要 copy 模組來 deepcopy 權重

def cnn_model(X_train, Y_train, Y_train_orig, X_test, Y_test_orig,
               learning_rate=0.05, num_epochs=10, batch_size=64,
               n_filters=8, filter_size=3,
               use_early_stopping=True, early_stopping_patience=5):

    (N_train, C_train, H_train, W_train) = X_train.shape
    num_classes = Y_train.shape[1]

    # 1. 初始化權重
    parameters = initialize_parameters(input_dim=(C_train, H_train, W_train),
                                       n_filters=n_filters,
                                       filter_size=filter_size,
                                       n_y=num_classes)

    # 儲存每個 epoch 的 loss 和 accuracy，用於後續繪圖
    costs = []
    accuracies = []

    # ≡≡≡ 早停機制初始化 ≡≡≡
```

```

if use_early_stopping:
    patience = early_stopping_patience
    epochs_no_improve = 0
    best_accuracy = 0.0
    # 用來儲存性能最好時的權重
    best_parameters = {}

# 2. 訓練迴圈
for i in range(num_epochs):
    epoch_start_time = time.time()
    epoch_cost = 0.0

    # 在每個 epoch 開始前，將資料隨機打亂
    permutation = cp.random.permutation(N_train)
    shuffled_X = X_train[permutation, :, :]
    shuffled_Y = Y_train[permutation, :]

    # Mini-batch 處理
    num_minibatches = N_train // batch_size

    for j in range(num_minibatches):
        start = j * batch_size
        end = start + batch_size
        minibatch_X = shuffled_X[start:end, :, :]
        minibatch_Y = shuffled_Y[start:end, :]

        # a. 前向傳播
        A_out, cache = forward_propagation(minibatch_X, parameters)

        # b. 計算損失
        cost = compute_loss(A_out, minibatch_Y)
        epoch_cost += cost

        # c. 反向傳播
        grads = backward_propagation(A_out, minibatch_Y, cache,
parameters)

        # d. 更新參數
        parameters = update_parameters(parameters, grads,
learning_rate)

    # 計算並記錄這個 epoch 的平均損失
    avg_epoch_cost = epoch_cost / num_minibatches
    costs.append(avg_epoch_cost)

    # 在每個 epoch 結束後，用測試集評估準確率

```

```

        predictions = predict(X_test, parameters)
        current_accuracy = calculate_accuracy(predictions, Y_test_orig)
        accuracies.append(current_accuracy)

    epoch_end_time = time.time()

    print(f"Epoch {i + 1}/{num_epochs} - Cost: {avg_epoch_cost:.6f} - Accuracy: {current_accuracy:.4f} - Time: {epoch_end_time - epoch_start_time:.2f}s")

    # 早停機制
    if use_early_stopping:
        if current_accuracy > best_accuracy:
            best_accuracy = current_accuracy
            epochs_no_improve = 0
            # deepcopy 目前最好的權重
            best_parameters = copy.deepcopy(parameters)
            print(f" -> Accuracy improved to {best_accuracy:.4f}! Saving model.")
        else:
            epochs_no_improve += 1
            print(f" -> Accuracy did not improve for {epochs_no_improve} epoch(s).")

        if epochs_no_improve >= patience:
            print(f"\nEarly stopping triggered after {patience} epochs with no improvement.")
            # 訓練結束，回傳歷史最佳的權重
            parameters = best_parameters
            break

    history = {"costs": costs, "accuracies": accuracies}
    return parameters, history

```

```

[ ]: # --- 預測與準確率計算的輔助函式 ---
def predict(X, parameters):
    A_out, _ = forward_propagation(X, parameters)
    # 找出每個樣本中，機率最高的那個類別的索引 (0-9)
    predictions = cp.argmax(A_out, axis=1)
    return predictions

def calculate_accuracy(predictions, labels):
    # 比較預測結果和正解標籤
    return cp.mean(predictions == labels)

```

## 5.2 Step 3-2: 訓練模型

```
[ ]: # 設定超參數
LEARNING_RATE = 0.05
NUM_EPOCHS = 100
BATCH_SIZE = 64
N_FILTERS = 32
FILTER_SIZE = 3
PATIENCE = 10

# 開始訓練
trained_parameters, history = cnn_model(
    train_x, train_y, train_y_orig,
    test_x, test_y_orig,
    learning_rate=LEARNING_RATE,
    num_epochs=NUM_EPOCHS,
    batch_size=BATCH_SIZE,
    n_filters=N_FILTERS,
    filter_size=FILTER_SIZE,
    use_early_stopping=True,
    # 如果連續 PATIENCE 次 accuracy 沒有提升就停止
    early_stopping_patience=PATIENCE
)
```

```
Epoch 1/100 - Cost: 0.402958 - Accuracy: 0.9508 - Time: 19.36s
-> Accuracy improved to 0.9508! Saving model.
Epoch 2/100 - Cost: 0.134227 - Accuracy: 0.9645 - Time: 8.23s
-> Accuracy improved to 0.9645! Saving model.
Epoch 3/100 - Cost: 0.098766 - Accuracy: 0.9617 - Time: 8.31s
-> Accuracy did not improve for 1 epoch(s).
Epoch 4/100 - Cost: 0.076742 - Accuracy: 0.9744 - Time: 8.39s
-> Accuracy improved to 0.9744! Saving model.
Epoch 5/100 - Cost: 0.058125 - Accuracy: 0.9736 - Time: 8.42s
-> Accuracy did not improve for 1 epoch(s).
Epoch 6/100 - Cost: 0.046284 - Accuracy: 0.9814 - Time: 8.45s
-> Accuracy improved to 0.9814! Saving model.
Epoch 7/100 - Cost: 0.038609 - Accuracy: 0.9812 - Time: 8.55s
-> Accuracy did not improve for 1 epoch(s).
Epoch 8/100 - Cost: 0.031190 - Accuracy: 0.9822 - Time: 8.56s
-> Accuracy improved to 0.9822! Saving model.
Epoch 9/100 - Cost: 0.024794 - Accuracy: 0.9818 - Time: 8.52s
-> Accuracy did not improve for 1 epoch(s).
Epoch 10/100 - Cost: 0.020114 - Accuracy: 0.9840 - Time: 8.47s
-> Accuracy improved to 0.9840! Saving model.
Epoch 11/100 - Cost: 0.016269 - Accuracy: 0.9814 - Time: 8.94s
-> Accuracy did not improve for 1 epoch(s).
Epoch 12/100 - Cost: 0.014241 - Accuracy: 0.9848 - Time: 8.44s
-> Accuracy improved to 0.9848! Saving model.
```

Epoch 13/100 - Cost: 0.010336 - Accuracy: 0.9839 - Time: 8.39s  
-> Accuracy did not improve for 1 epoch(s).  
Epoch 14/100 - Cost: 0.008706 - Accuracy: 0.9843 - Time: 8.44s  
-> Accuracy did not improve for 2 epoch(s).  
Epoch 15/100 - Cost: 0.005845 - Accuracy: 0.9854 - Time: 8.45s  
-> Accuracy improved to 0.9854! Saving model.  
Epoch 16/100 - Cost: 0.004464 - Accuracy: 0.9826 - Time: 8.45s  
-> Accuracy did not improve for 1 epoch(s).  
Epoch 17/100 - Cost: 0.006639 - Accuracy: 0.9852 - Time: 8.43s  
-> Accuracy did not improve for 2 epoch(s).  
Epoch 18/100 - Cost: 0.004916 - Accuracy: 0.9826 - Time: 8.48s  
-> Accuracy did not improve for 3 epoch(s).  
Epoch 19/100 - Cost: 0.003173 - Accuracy: 0.9854 - Time: 8.46s  
-> Accuracy did not improve for 4 epoch(s).  
Epoch 20/100 - Cost: 0.001940 - Accuracy: 0.9843 - Time: 8.43s  
-> Accuracy did not improve for 5 epoch(s).  
Epoch 21/100 - Cost: 0.001605 - Accuracy: 0.9855 - Time: 8.46s  
-> Accuracy improved to 0.9855! Saving model.  
Epoch 22/100 - Cost: 0.000699 - Accuracy: 0.9855 - Time: 9.67s  
-> Accuracy did not improve for 1 epoch(s).  
Epoch 23/100 - Cost: 0.000546 - Accuracy: 0.9853 - Time: 9.06s  
-> Accuracy did not improve for 2 epoch(s).  
Epoch 24/100 - Cost: 0.000499 - Accuracy: 0.9859 - Time: 8.41s  
-> Accuracy improved to 0.9859! Saving model.  
Epoch 25/100 - Cost: 0.000452 - Accuracy: 0.9850 - Time: 8.43s  
-> Accuracy did not improve for 1 epoch(s).  
Epoch 26/100 - Cost: 0.000570 - Accuracy: 0.9856 - Time: 8.44s  
-> Accuracy did not improve for 2 epoch(s).  
Epoch 27/100 - Cost: 0.000286 - Accuracy: 0.9859 - Time: 8.50s  
-> Accuracy did not improve for 3 epoch(s).  
Epoch 28/100 - Cost: 0.000190 - Accuracy: 0.9860 - Time: 8.41s  
-> Accuracy improved to 0.9860! Saving model.  
Epoch 29/100 - Cost: 0.000160 - Accuracy: 0.9857 - Time: 8.70s  
-> Accuracy did not improve for 1 epoch(s).  
Epoch 30/100 - Cost: 0.000152 - Accuracy: 0.9854 - Time: 8.44s  
-> Accuracy did not improve for 2 epoch(s).  
Epoch 31/100 - Cost: 0.000148 - Accuracy: 0.9856 - Time: 8.46s  
-> Accuracy did not improve for 3 epoch(s).  
Epoch 32/100 - Cost: 0.000131 - Accuracy: 0.9860 - Time: 8.42s  
-> Accuracy did not improve for 4 epoch(s).  
Epoch 33/100 - Cost: 0.000127 - Accuracy: 0.9860 - Time: 8.48s  
-> Accuracy did not improve for 5 epoch(s).  
Epoch 34/100 - Cost: 0.000113 - Accuracy: 0.9856 - Time: 8.53s  
-> Accuracy did not improve for 6 epoch(s).  
Epoch 35/100 - Cost: 0.000103 - Accuracy: 0.9858 - Time: 8.43s  
-> Accuracy did not improve for 7 epoch(s).  
Epoch 36/100 - Cost: 0.000099 - Accuracy: 0.9857 - Time: 8.49s  
-> Accuracy did not improve for 8 epoch(s).



Epoch 37/100 - Cost: 0.000095 - Accuracy: 0.9856 - Time: 8.48s  
-> Accuracy did not improve for 9 epoch(s).  
Epoch 38/100 - Cost: 0.000090 - Accuracy: 0.9854 - Time: 8.46s  
-> Accuracy did not improve for 10 epoch(s).

Early stopping triggered after 10 epochs with no improvement.

### 5.3 Step 3-3: 儲存及載入參數的函數

儲存模型的權重參數

參數:

- parameters: 包含 CuPy 陣列的權重字典
- filename: 要儲存的檔案名稱

```
[ ]: def save_parameters(parameters, filename="my_cnn_model.npz"):
    print(f"Saving parameters to {filename}...")

    # 建立一個新的字典，用來存放轉換後的 NumPy 陣列
    numpy_params = {}
    for key, value in parameters.items():
        # 使用 cp.asnumpy() 將 CuPy 陣列轉為 NumPy 陣列
        numpy_params[key] = cp.asnumpy(value)

    # 使用 np.savez_compressed 儲存，**numpy_params 會自動解包成關鍵字參數
    np.savez_compressed(filename, **numpy_params)

    print("Parameters saved successfully.")
```

```
[2]: import os

def load_parameters(filename="my_cnn_model.npz"):
    if not os.path.exists(filename):
        print(f"Error: File '{filename}' not found.")
        return None

    print(f"Loading parameters from {filename}...")

    # 載入 .npz 檔案，它是一個類字典物件
    npzfile = np.load(filename)

    # 建立一個新的字典，用來存放轉換後的 CuPy 陣列
    parameters = {}
    for key in npzfile.keys():
        # 使用 cp.asarray() 將 NumPy 陣列轉為 CuPy 陣列
        parameters[key] = cp.asarray(npzfile[key])

    print("Parameters loaded successfully.")
```

```
return parameters
```

```
[ ]: save_parameters(trained_parameters, "my_cnn_model.npz")
```

Saving parameters to my\_cnn\_model.npz...  
Parameters saved successfully.

## 6 Step 4: 最終評估與報告

### 6.1 Step 4-1: 最終預測

```
[ ]: trained_parameters = load_parameters("my_cnn_model.npz")
```

Loading parameters from my\_cnn\_model.npz...  
Parameters loaded successfully.

```
[ ]: # 1. 進行最終預測
print("Making final predictions on the test set...")
test_predictions = predict(test_x, trained_parameters)

# 2. 確認最終準確率
final_accuracy = calculate_accuracy(test_predictions, test_y_orig)
# 使用 cp.asnumpy() 將 CuPy 純量轉為 Python 純量以便格式化輸出
print(f"Final model accuracy: {cp.asnumpy(final_accuracy)*100:.2f}%")
```

Making final predictions on the test set...  
Final model accuracy: 98.60%

```
[ ]: # 3. 建立混淆矩陣
num_classes = 10
# 初始化一個 10x10 的 CuPy 矩陣
confusion_matrix_cp = cp.zeros((num_classes, num_classes), dtype=int)

# 遍歷所有測試集樣本，填充混淆矩陣
for i in range(len(test_y_orig)):
    true_label = test_y_orig[i]
    predicted_label = test_predictions[i]
    confusion_matrix_cp[true_label, predicted_label] += 1
```

```
[ ]: # 4. 視覺化混淆矩陣
# 記得要先用 .asnumpy() 將矩陣從 GPU 移回 CPU
confusion_matrix_np = cp.asnumpy(confusion_matrix_cp)


import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 8))
sns.heatmap(
```

```

confusion_matrix_np,
annot=True,      # 在格子中顯示數字
fmt="d",         # 將數字格式化為整數
cmap="Blues",    # 使用藍色系的色盤
xticklabels=range(num_classes),
yticklabels=range(num_classes),
)
plt.title("Confusion Matrix Heatmap")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```



## 6.2 Step 4-2: 計算精確率 (Precision) 與召回率 (Recall)

```
[ ]: # 我們直接在 CuPy 陣列上進行計算
print("--- Performance Metrics per Class ---")
for i in range(num_classes):
    true_positives = confusion_matrix_cp[i, i]
    # 該 Col 的總和 (所有被預測為 i 的) - TP
    false_positives = cp.sum(confusion_matrix_cp[:, i]) - \
true_positives
    # 該 Row 的總和 (所有實際上是 i 的) - TP
    false_negatives = cp.sum(confusion_matrix_cp[i, :]) - \
true_positives

    # 計算 Precision 和 Recall
    precision = true_positives / (true_positives + false_positives)
    recall = true_positives / (true_positives + false_negatives)

    print(f"數字 '{i}':")
    # 使用 .asnumpy() 以便格式化輸出
    print(f" 精確率 (Precision): {cp.asnumpy(precision):.4f}")
    print(f" 召回率 (Recall): {cp.asnumpy(recall):.4f}")
```

--- Performance Metrics per Class ---

數字 '0':

精確率 (Precision): 0.9868

召回率 (Recall): 0.9918

數字 '1':

精確率 (Precision): 0.9956

召回率 (Recall): 0.9947

數字 '2':

精確率 (Precision): 0.9855

召回率 (Recall): 0.9874

數字 '3':

精確率 (Precision): 0.9833

召回率 (Recall): 0.9921

數字 '4':

精確率 (Precision): 0.9857

召回率 (Recall): 0.9796

數字 '5':

精確率 (Precision): 0.9854

召回率 (Recall): 0.9865

數字 '6':

精確率 (Precision): 0.9854

召回率 (Recall): 0.9854

數字 '7':

精確率 (Precision): 0.9825

召回率 (Recall): 0.9854

數字 '8':

精確率 (Precision): 0.9906  
召回率 (Recall): 0.9764  
數字 '9':  
精確率 (Precision): 0.9782  
召回率 (Recall): 0.9792

### 6.3 Step 4-3: 繪制曲線圖

```
[ ]: # 將 CuPy 陣列轉為 NumPy 陣列
# 1. 從 history 中取出資料
costs_cp = cp.array(history['costs'])
accuracies_cp = cp.array(history['accuracies'])

# 2. 使用 asnumpy() 轉換
costs_np = cp.asnumpy(costs_cp)
accuracies_np = cp.asnumpy(accuracies_cp)

# 開始繪圖
# 建立一個 1x2 的子圖畫布，方便同時顯示兩張圖
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))

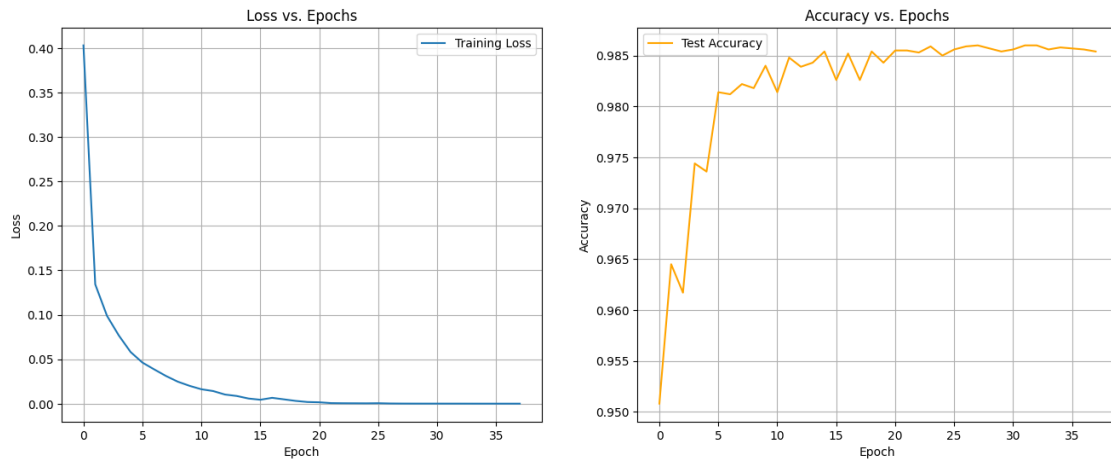
# 設定一個主標題
fig.suptitle('CNN Model Learning Curves', fontsize=16)

# 圖一：損失 (Loss) vs. 世代 (Epochs)
ax1.plot(costs_np, label='Training Loss')
ax1.set_title("Loss vs. Epochs")
ax1.set_xlabel("Epoch")
ax1.set_ylabel("Loss")
ax1.grid(True)
ax1.legend()

# 圖二：準確率 (Accuracy) vs. 世代 (Epochs)
ax2.plot(accuracies_np, label='Test Accuracy', color='orange')
ax2.set_title("Accuracy vs. Epochs")
ax2.set_xlabel("Epoch")
ax2.set_ylabel("Accuracy")
ax2.grid(True)
ax2.legend()

# 顯示圖表
plt.show()
```

CNN Model Learning Curves



#### 6.4 Step 4-4: 加分：視覺化卷積核 (Visualize Filters)

```
[ ]: trained_parameters = load_parameters("my_cnn_model.npz")
assert trained_parameters is not None
if 'trained_parameters' in locals() or 'trained_parameters' in_
globals():
    print("讀取已訓練的 MNIST 模型權重...")

# 1. 從權重字典中取出第一個卷積層的權重 (W_conv)
# 它的 shape 是 (n_filters, C_in, f, f)，在此為 (32, 1, 3, 3)
W_conv_cp = trained_parameters['W_conv']

# 2. 將 CuPy 陣列移至 CPU 並轉為 NumPy 陣列，以便 Matplotlib 處理
W_conv_np = cp.asnumpy(W_conv_cp)

# 3. 設定視覺化圖表
# 我們有 32 個 filters，可以用一個 4x8 的網格來顯示
num_filters = W_conv_np.shape[0]
fig, axes = plt.subplots(4, 8, figsize=(12, 6))

# 設定主標題
fig.suptitle('Visualization of Learned Convolutional Filters_
(MNIST Model)', fontsize=16)

# 4. 遍歷所有 filters 並在子圖中繪製
for i, ax in enumerate(axes.flat):
    # 檢查是否還有 filter 可以顯示
    if i < num_filters:
        # 取出第 i 個 filter 的權重。
        # 由於 MNIST 是灰階，輸入頻道只有 1，所以我們用 [i, 0, :, :]
```

```

filt = W_conv_np[i, 0, :, :]

# 使用 imshow 繪製 filter, 'gray' 色彩映射適合顯示權重
ax.imshow(filt, cmap='gray')

# 隱藏座標軸刻度，讓圖表更簡潔
ax.set_xticks([])
ax.set_yticks([])

# 調整子圖之間的間距
plt.tight_layout()
# 顯示圖表
plt.show()


else:
    print("錯誤：找不到 'trained_parameters'。請先確定您已成功訓練並載入_
    ↪ MNIST 模型。")

```

Loading parameters from my\_cnn\_model.npz...

Parameters loaded successfully.

讀取已訓練的 MNIST 模型權重...



## 7 Bonus: CIFAR-10

```

[ ]: import gc # 引入垃圾回收模組

# 刪除所有不再需要的 MNIST 變數
mnist_vars_to_delete = [

```

```

    'train_x', 'train_y', 'train_x_orig', 'train_y_orig',
    'test_x', 'test_y', 'test_x_orig', 'test_y_orig',
    'trained_parameters',
    'test_predictions',
    'confusion_matrix_cp'
]

for var_name in mnist_vars_to_delete:
    if var_name in locals() or var_name in globals():
        del globals()[var_name]
        print(f"變數 '{var_name}' 已刪除。")

# 強制進行垃圾回收
gc.collect()

# 清空 CuPy 的記憶體池
mempool = cp.get_default_memory_pool()
mempool.free_all_blocks()

print("\nCuPy 記憶體池已清空。")
print(f"目前 VRAM 已分配 (Used Bytes): {mempool.used_bytes()} bytes")
print(f"目前 VRAM 總量 (Total Bytes): {mempool.total_bytes()} bytes")

```

變數 'train\_x' 已刪除。  
 變數 'train\_y' 已刪除。  
 變數 'train\_x\_orig' 已刪除。  
 變數 'train\_y\_orig' 已刪除。  
 變數 'test\_x' 已刪除。  
 變數 'test\_y' 已刪除。  
 變數 'test\_x\_orig' 已刪除。  
 變數 'test\_y\_orig' 已刪除。  
 變數 'trained\_parameters' 已刪除。  
 變數 'test\_predictions' 已刪除。  
 變數 'confusion\_matrix\_cp' 已刪除。

CuPy 記憶體池已清空。  
 目前 VRAM 已分配 (Used Bytes): 1670149120 bytes  
 目前 VRAM 總量 (Total Bytes): 3036867584 bytes

## 7.1 Bonus Step 1: 下載並加載 CIFAR-10

```
[ ]: !wget https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
```

```

--2025-06-21 15:19:58--  https://www.cs.toronto.edu/~kriz/
    ↪ cifar-10-python.tar.gz
Resolving www.cs.toronto.edu (www.cs.toronto.edu)... 128.100.3.30
Connecting to www.cs.toronto.edu (www.cs.toronto.edu)|128.100.3.30|:443...
connected.

```



```
HTTP request sent, awaiting response... 200 OK
Length: 170498071 (163M) [application/x-gzip]
Saving to: 'cifar-10-python.tar.gz'
```

```
cifar-10-python.tar 100%[=====>] 162.60M  1.05MB/s   in_
  2m 53s
```

```
2025-06-21 15:22:53 (961 KB/s) - 'cifar-10-python.tar.gz' saved
[170498071/170498071]
```

```
[ ]: !tar zxvf cifar-10-python.tar.gz
```

```
cifar-10-batches-py/
cifar-10-batches-py/data_batch_4
cifar-10-batches-py/readme.html
cifar-10-batches-py/test_batch
cifar-10-batches-py/data_batch_3
cifar-10-batches-py/batches.meta
cifar-10-batches-py/data_batch_2
cifar-10-batches-py/data_batch_5
cifar-10-batches-py/data_batch_1
```

```
[ ]: import pickle
import os
```

```
[ ]: def load_cifar10(root):
    def load_batch(filename):
        with open(os.path.join(root, filename), 'rb') as f:
            # 使用 latin1 編碼來讀取 Python 2 生成的 pickle 檔案
            datadict = pickle.load(f, encoding='latin1')
            X = datadict['data']
            Y = datadict['labels']
            # 將資料塑形為 (樣本數, 色彩頻道, 高度, 寬度)
            # 原始資料格式為 (N, 3072), 其中 3072 = 3 * 32 * 32
            # 儲存順序是 RRR...GGG...BBB...
            X = X.reshape(10000, 3, 32, 32)
            Y = np.array(Y)
            return X, Y

    # 載入訓練資料 (共 5 個 batch)
    xs, ys = [], []
    for i in range(1, 6):
        x, y = load_batch(f"data_batch_{i}")
        xs.append(x)
        ys.append(y)

    train_x = np.concatenate(xs)
```

```

train_y = np.concatenate(ys)

# 載入測試資料
test_x, test_y = load_batch("test_batch")

# 載入類別名稱
with open(os.path.join(root, 'batches.meta'), 'rb') as f:
    meta = pickle.load(f, encoding='latin1')
    class_names = meta['label_names']

return train_x, train_y, test_x, test_y, class_names

```

```

[ ]: cifar10_train_x_np, cifar10_train_y_np, cifar10_test_x_np, \
      ↪ cifar10_test_y_np, cifar10_class_names = \
      ↪ load_cifar10("cifar-10-batches-py")
print("\n--- CIFAR-10 資料維度 (NumPy) ---")
print("訓練資料 X:", cifar10_train_x_np.shape)
print("訓練標籤 Y:", cifar10_train_y_np.shape)
print("測試資料 X:", cifar10_test_x_np.shape)
print("測試標籤 Y:", cifar10_test_y_np.shape)
print("類別名稱:", cifar10_class_names)

```

--- CIFAR-10 資料維度 (NumPy) ---

訓練資料 X: (50000, 3, 32, 32)

訓練標籤 Y: (50000,)

測試資料 X: (10000, 3, 32, 32)

測試標籤 Y: (10000,)

類別名稱: ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', \
 ↪ 'frog', 'horse',  
'ship', 'truck']

## 7.2 Bonus Step 2: 轉為 CuPy

```

[ ]: # 將 NumPy 陣列轉換為 CuPy 陣列
cifar10_train_x_orig = cp.asarray(cifar10_train_x_np)
cifar10_train_y_orig = cp.asarray(cifar10_train_y_np)
cifar10_test_x_orig = cp.asarray(cifar10_test_x_np)
cifar10_test_y_orig = cp.asarray(cifar10_test_y_np)

# --- 正規化 (Normalize) ---
# 將像素值從 0-255 正規化到 0-1 之間
cifar10_train_x = cifar10_train_x_orig / 255.0
cifar10_test_x = cifar10_test_x_orig / 255.0

# --- 對 Label 進行 One-Hot Encode ---
num_classes_cifar10 = 10

```

```

cifar10_train_y = one_hot_encode(cifar10_train_y_orig,
    ↪ num_classes_cifar10)
cifar10_test_y = one_hot_encode(cifar10_test_y_orig,
    ↪ num_classes_cifar10)

print("\n--- CIFAR-10 資料維度 (CuPy) ---")
print("train_x shape:", cifar10_train_x.shape)
print("train_y shape:", cifar10_train_y.shape)
print("test_x shape:", cifar10_test_x.shape)
print("test_y shape:", cifar10_test_y.shape)
print("Data type:", type(cifar10_train_x))

```

```

--- CIFAR-10 資料維度 (CuPy) ---
train_x shape: (50000, 3, 32, 32)
train_y shape: (50000, 10)
test_x shape: (10000, 3, 32, 32)
test_y shape: (10000, 10)
Data type: <class 'cupy.ndarray'>

```

### 7.3 Bonus Step 3: 訓練

```

[ ]: # 設定 hyperparameters
LEARNING_RATE_CIFAR = 0.05
NUM_EPOCHS_CIFAR = 150 # 增加 Epoch 上限
BATCH_SIZE_CIFAR = 128
N_FILTERS_CIFAR = 32 # 增加 filter 數量以捕捉更複雜的特徵
FILTER_SIZE_CIFAR = 3
PATIENCE_CIFAR = 15

# 開始訓練
cifar10_params, cifar10_history = cnn_model(
    cifar10_train_x, cifar10_train_y, cifar10_train_y_orig,
    cifar10_test_x, cifar10_test_y_orig,
    learning_rate=LEARNING_RATE_CIFAR,
    num_epochs=NUM_EPOCHS_CIFAR,
    batch_size=BATCH_SIZE_CIFAR,
    n_filters=N_FILTERS_CIFAR,
    filter_size=FILTER_SIZE_CIFAR,
    use_early_stopping=True,
    early_stopping_patience=PATIENCE_CIFAR
)

# 儲存訓練好的 CIFAR-10 模型權重
save_parameters(cifar10_params, "my_cifar10_cnn_model.npz")

```

```

Epoch 1/150 - Cost: 2.072176 - Accuracy: 0.3469 - Time: 10.08s
-> Accuracy improved to 0.3469! Saving model.

```

Epoch 2/150 - Cost: 1.753204 - Accuracy: 0.4441 - Time: 9.98s  
-> Accuracy improved to 0.4441! Saving model.  
Epoch 3/150 - Cost: 1.541613 - Accuracy: 0.4770 - Time: 9.69s  
-> Accuracy improved to 0.4770! Saving model.  
Epoch 4/150 - Cost: 1.422470 - Accuracy: 0.5105 - Time: 9.27s  
-> Accuracy improved to 0.5105! Saving model.  
Epoch 5/150 - Cost: 1.330337 - Accuracy: 0.5092 - Time: 9.34s  
-> Accuracy did not improve for 1 epoch(s).  
Epoch 6/150 - Cost: 1.248691 - Accuracy: 0.5237 - Time: 9.18s  
-> Accuracy improved to 0.5237! Saving model.  
Epoch 7/150 - Cost: 1.158501 - Accuracy: 0.5591 - Time: 9.20s  
-> Accuracy improved to 0.5591! Saving model.  
Epoch 8/150 - Cost: 1.072573 - Accuracy: 0.5901 - Time: 9.22s  
-> Accuracy improved to 0.5901! Saving model.  
Epoch 9/150 - Cost: 1.013198 - Accuracy: 0.5719 - Time: 9.25s  
-> Accuracy did not improve for 1 epoch(s).  
Epoch 10/150 - Cost: 0.944731 - Accuracy: 0.5490 - Time: 9.27s  
-> Accuracy did not improve for 2 epoch(s).  
Epoch 11/150 - Cost: 0.883016 - Accuracy: 0.5827 - Time: 9.27s  
-> Accuracy did not improve for 3 epoch(s).  
Epoch 12/150 - Cost: 0.830277 - Accuracy: 0.5769 - Time: 9.28s  
-> Accuracy did not improve for 4 epoch(s).  
Epoch 13/150 - Cost: 0.776206 - Accuracy: 0.5644 - Time: 9.26s  
-> Accuracy did not improve for 5 epoch(s).  
Epoch 14/150 - Cost: 0.728672 - Accuracy: 0.5807 - Time: 9.26s  
-> Accuracy did not improve for 6 epoch(s).  
Epoch 15/150 - Cost: 0.684573 - Accuracy: 0.5901 - Time: 9.24s  
-> Accuracy did not improve for 7 epoch(s).  
Epoch 16/150 - Cost: 0.624721 - Accuracy: 0.5801 - Time: 9.23s  
-> Accuracy did not improve for 8 epoch(s).  
Epoch 17/150 - Cost: 0.587469 - Accuracy: 0.5757 - Time: 9.23s  
-> Accuracy did not improve for 9 epoch(s).  
Epoch 18/150 - Cost: 0.542030 - Accuracy: 0.5882 - Time: 9.25s  
-> Accuracy did not improve for 10 epoch(s).  
Epoch 19/150 - Cost: 0.508496 - Accuracy: 0.5852 - Time: 9.34s  
-> Accuracy did not improve for 11 epoch(s).  
Epoch 20/150 - Cost: 0.482322 - Accuracy: 0.5689 - Time: 9.25s  
-> Accuracy did not improve for 12 epoch(s).  
Epoch 21/150 - Cost: 0.442902 - Accuracy: 0.5829 - Time: 9.27s  
-> Accuracy did not improve for 13 epoch(s).  
Epoch 22/150 - Cost: 0.412536 - Accuracy: 0.5788 - Time: 9.26s  
-> Accuracy did not improve for 14 epoch(s).  
Epoch 23/150 - Cost: 0.372546 - Accuracy: 0.5764 - Time: 10.03s  
-> Accuracy did not improve for 15 epoch(s).

Early stopping triggered after 15 epochs with no improvement.  
Saving parameters to my\_cifar10\_cnn\_model.npz...  
Parameters saved successfully.

## 7.4 Bonus Step 4: 效能評估

```
[ ]: cifar10_trained_parameters = load_parameters("my_cifar10_cnn_model.
↳npz")
```

Loading parameters from my\_cifar10\_cnn\_model.npz...  
Parameters loaded successfully.

```
[ ]: # 1. 進行最終預測
print("在 CIFAR-10 測試集上進行最終預測")
cifar10_test_predictions = predict(cifar10_test_x,
↳cifar10_trained_parameters)

# 2. 確認最終準確率
final_accuracy_cifar10 = calculate_accuracy(cifar10_test_predictions,
↳cifar10_test_y_orig)
print(f"CIFAR-10 模型最終準確率: {cp.asnumpy(final_accuracy_cifar10)*100:.
↳2f}%")
```

在 CIFAR-10 測試集上進行最終預測  
CIFAR-10 模型最終準確率: 59.01%


```
[ ]: # 3. 建立混淆矩陣
confusion_matrix_cifar10_cp = cp.zeros((num_classes_cifar10,
↳num_classes_cifar10), dtype=int)
for i in range(len(cifar10_test_y_orig)):
    true_label = cifar10_test_y_orig[i]
    predicted_label = cifar10_test_predictions[i]
    confusion_matrix_cifar10_cp[true_label, predicted_label] += 1

# 4. 視覺化混淆矩陣 (使用類別名稱)
confusion_matrix_cifar10_np = cp.asnumpy(confusion_matrix_cifar10_cp)

import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(12, 10))
sns.heatmap(
    confusion_matrix_cifar10_np,
    annot=True,
    fmt="d",
    cmap="Blues",
    xticklabels=cifar10_class_names,
    yticklabels=cifar10_class_names,
)
plt.title("CIFAR-10 Confusion Matrix Heatmap")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
```

```
plt.show()
```



```
[ ]: # 5. 計算精確率與召回率
print("\n--- CIFAR-10 各類別的性能指標 ---")
for i in range(num_classes_cifar10):
    true_positives = confusion_matrix_cifar10_cp[i, i]
    false_positives = cp.sum(confusion_matrix_cifar10_cp[:, i]) - \
true_positives
    false_negatives = cp.sum(confusion_matrix_cifar10_cp[i, :]) - \
true_positives

    precision = true_positives / (true_positives + false_positives + \
1e-9) # 避免除以零
    recall = true_positives / (true_positives + false_negatives + 1e-9)
```

```

print(f"類別 '{cifar10_class_names[i]}':")
print(f"  精確率 (Precision): {cp.asnumpy(precision):.4f}")
print(f"  召回率 (Recall):    {cp.asnumpy(recall):.4f}")

```

--- CIFAR-10 各類別的性能指標 ---

```

類別 'airplane':
  精確率 (Precision): 0.6383
  召回率 (Recall):    0.6440
類別 'automobile':
  精確率 (Precision): 0.7225
  召回率 (Recall):    0.6950
類別 'bird':
  精確率 (Precision): 0.5250
  召回率 (Recall):    0.4200
類別 'cat':
  精確率 (Precision): 0.4277
  召回率 (Recall):    0.3430
類別 'deer':
  精確率 (Precision): 0.5254
  召回率 (Recall):    0.4870
類別 'dog':
  精確率 (Precision): 0.4930
  召回率 (Recall):    0.5280
類別 'frog':
  精確率 (Precision): 0.5737
  召回率 (Recall):    0.6770
類別 'horse':
  精確率 (Precision): 0.6308
  召回率 (Recall):    0.6850
類別 'ship':
  精確率 (Precision): 0.6854
  召回率 (Recall):    0.7430
類別 'truck':
  精確率 (Precision): 0.6293
  召回率 (Recall):    0.6790

```

[ ]: # 6. 繪製學習曲線

```

costs_cifar10_cp = cp.array(cifar10_history['costs'])
accuracies_cifar10_cp = cp.array(cifar10_history['accuracies'])
costs_cifar10_np = cp.asnumpy(costs_cifar10_cp)
accuracies_cifar10_np = cp.asnumpy(accuracies_cifar10_cp)

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))
fig.suptitle('CIFAR-10 Model Learning Curves', fontsize=16)

ax1.plot(costs_cifar10_np, label='Training Loss')

```


```

ax1.set_title("Loss vs. Epochs")
ax1.set_xlabel("Epoch")
ax1.set_ylabel("Loss")
ax1.grid(True)
ax1.legend()

ax2.plot(accuracies_cifar10_np, label='Test Accuracy', color='orange')
ax2.set_title("Accuracy vs. Epochs")
ax2.set_xlabel("Epoch")
ax2.set_ylabel("Accuracy")
ax2.grid(True)
ax2.legend()

plt.show()

```



## 7.5 Bonus Step 5: 加分題：視覺化卷積核

```

[7]: import numpy as np
import cupy as cp
import matplotlib.pyplot as plt

```

```

[11]: trained_parameters = load_parameters("my_cifar10_cnn_model.npz")
assert trained_parameters is not None
if 'trained_parameters' in locals() or 'trained_parameters' in_
globals():
    print("讀取已訓練的 CIFAR-10 模型權重...")

    # 1. 從權重字典中取出第一個卷積層的權重 (W_conv)
    W_conv_cp = trained_parameters['W_conv']

```



```

# 2. 將 CuPy 陣列移至 CPU 並轉為 NumPy 陣列，以便 Matplotlib 處理
W_conv_np = cp.asnumpy(W_conv_cp)

# 3. 設定視覺化圖表
# 我們有 32 個 filters，可以用一個 4x8 的網格來顯示
num_filters = W_conv_np.shape[0]
fig, axes = plt.subplots(4, 8, figsize=(12, 6))

# 設定主標題
fig.suptitle('Visualization of Learned Convolutional Filters_
↳(CIFAR-10)', fontsize=16)

# 4. 遍歷所有 filters 並在子圖中繪製
for i, ax in enumerate(axes.flat):
    # 檢查是否還有 filter 可以顯示
    if i < num_filters:
        # 取出第 i 個 filter 的權重。
        filt = W_conv_np[i, :, :, :]

        f_min, f_max = filt.min(), filt.max()
        filt_normalized = (filt - f_min) / (f_max - f_min)

        # 轉置維度從 (C, H, W) 到 (H, W, C) 以符合 imshow 的格式
        filt_rgb = filt_normalized.transpose(1, 2, 0)

        # 使用 imshow 繪製 filter
        ax.imshow(filt_rgb)

        # 隱藏座標軸刻度，讓圖表更簡潔
        ax.set_xticks([])
        ax.set_yticks([])

    # 調整子圖之間的間距
    plt.tight_layout()
    # 顯示圖表
    plt.show()

else:
    print("錯誤：找不到 'trained_parameters'。請先確定您已成功訓練並載入_
↳MNIST 模型。")


```

Loading parameters from my\_cifar10\_cnn\_model.npz...

Parameters loaded successfully.

讀取已訓練的 CIFAR-10 模型權重...

Visualization of Learned Convolutional Filters (CIFAR-10)



## 8 Conclusion

本次作業成功地從零開始建構了一個功能完整的卷積神經網路，不僅達成了作業 hw4.pdf 的所有基本要求，也完成了加分挑戰項目，從而驗證了對 CNN 核心概念的深入理解。

在 MNIST 資料集上，模型表現出色，達到 98.60% 的高準確率。從學習曲線圖中可以看到，訓練損失穩定下降，而測試準確率則快速上升後趨於平穩，顯示模型有效地學習到了數字特徵。透過視覺化卷積核，可以觀察到模型在第一層學會了辨識數字的基礎筆劃，如邊緣、角落與線條等特徵。此外，提早停止（Early Stopping）機制的導入，成功地在模型性能達到巔峰時結束訓練，有效防止了過擬合。

在 CIFAR-10 加分項目中，模型雖然能夠運作，但 59.01% 的準確率遠低於 MNIST 的成果。這符合預期，因為 CIFAR-10 是由更複雜的自然圖像組成，其特徵（如紋理、顏色、物體形狀）遠比 MNIST 的單色筆劃複雜。本次實作的淺層 CNN 結構對於捕捉如此複雜的特徵能力有限。儘管準確率不高，但成功將模型拓展至處理 3 通道的彩色圖像，並完成端到端的訓練與評估，本身即證明了模型的擴展性與彈性。

總體而言，本次專案最大的收穫在於透過親手打造每一個環節，深刻體會了 CNN 的數學原理與運作流程。從巢狀迴圈到 im2col 的效能最佳化過程，也讓人理解到向量化運算在深度學習中的關鍵性。而 CuPy 的使用，則展示了在不依賴高階框架的前提下，利用 GPU 大幅提升訓練效率的可行性。